

# **COMP9311 - Database Systems**

## **Week 1**

### **Introduction, Data Modeling, and ER Notation**

# **School of Computer Science and Engineering (CSE)**

**Lecturer: Dr. Rachid Hamadi**

**Email: rhamadi@cse.unsw.edu.au**

**Course Website:**

**<http://www.cse.unsw.edu.au/~cs9311>**

# Course Information

- Lectures: **Wednesdays 15:00 - 18:00**

Location: **ChemScM17**

- Labs: **Weeks 2 – 13, 2 hours per week in labs**

- Consultation: **Wednesdays 14:00 - 15:00**

Location: **K17-203**

# Course Information (cont'd)

- This semester, there will be three (3) assignments and a final exam
- Assignments are worth **40%**:
  - Assignment 1: ER Diagram / Relational Mapping (**10%**) (due **Week 5**)
  - Assignment 2: SQL Queries and Functions (15%) (due **Week 8**)
  - Assignment 3: Normalisation / Relational Algebra / Transaction (**15%**) (due **Week 11**)
- **Penalty for late submission:** 10% reduction of the total mark for each day late.

# Course Information (cont'd)

- **Final Exam:**

Two (2) hours **written** exam

Worth **60%** during exam period

- **Final Mark:**

FinalMark = A1Mark + A2Mark + A3Mark + ExamMark

- **Hurdle:**

To pass this course, you must obtain **at least 40%** in the final exam with a combined total of at least 50%

# Course Information (cont'd)

- Textbook:
  - [Fundamentals of Database Systems](#), Elmasri and Navathe, 7th edition, 2016, Addison-Wesley
- Reference books:
  - [Database System Concepts](#), Silberschatz, Korth, Sudarshan, 6th edition, 2010, McGraw-Hill
  - [Database Management Systems](#), Ramakrishnan and Gehrke, 3rd edition, 2003, McGraw-Hill
  - [Database Systems: The Complete Book](#), Garcia-Molina, Ullman, Widom, 2nd edition, 2008, Prentice-Hall
  - [Database Systems: An Application-Oriented Approach](#) Kifer, Bernstein, Lewis, 2nd edition (Complete Version), 2006, Addison-Wesley

Note: Earlier editions are fine for this course

# Course Schedule

The following is an **approximate** guide to the sequence of topics in this course.

Week	Lectures	Assignments
1	Introduction, Data Modelling, ER Notation	
2	Relational Model, ER-Relational Mapping, SQL Schemas	
3	DBMSs, Databases, Data Modification	
4	SQL Queries	
5	More SQL Queries, Stored Procedures, PLpgSQL	<b>Assignment 1 Due</b>
6	Extending SQL: Queries, Functions, Aggregates, Triggers	
7	More Triggers, Programming with Databases	
8	Catalogs, Privileges	<b>Assignment 2 Due</b>
9	Relational Design Theory, Normal Forms	
-	Non-teaching week (mid-semester break)	-
10	Relational Algebra, Query Processing	
11	Transaction Processing, Concurrency Control	<b>Assignment 3 Due</b>
12	Course Review	

# Databases in CSE

COMP9311 introduces foundations & technology of databases

- Skills: how to build database-backed applications
- Theory: how do you know that what you built was any good

After COMP9311 you can go on to study ...

- COMP9315: how to build relational DBMSs (write your own Oracle)
- COMP9318: techniques for data mining (discovering patterns in DB)
- COMP6714: information retrieval, web search (dealing with text data)
- COMP9319: web search and data compression (dealing searching compressed web data)
- COMP932(1|2|3): service-oriented computing, which relies on DB background



# Introduction

- Database Applications:
  - Banking System,
  - Stock Market,
  - Transportation,
  - Social Network,
  - Marine Data Analysis,
  - Criminal Analysis and Control,
  - Now, Big Data ...

# Introduction (cont'd)

**Intelligent Transportation**



**Business Services**



**Natural Disasters**



**Public Health**




**Modern Military**



**Tourism Development**

# More Applications



Chemical Compounds



Social Network



Collaboration Graph



Road Network

# Graphs Could be Big!



- 1.23 billion active users in 2013
- 117 billion friendships in 2013



- 645 million users in 2013
- 1.7 billion tweets/month in 2013

# Databases: Important Themes

The field of **databases** deals with:

- **Data** ... representing application scenarios
- **Relationships** ... amongst data items
- **Constraints** ... on data and relationships
- **Redundancy** ... one source for each data item
- **Data manipulation** ... declarative, procedural
- **Transactions** ... multiple actions, atomic effect
- **Concurrency** ... multiple users sharing data
- **Scale** ... massive amounts of data

# What is data?

- *Data* (Elmasri/Navathe)
  - Known facts that can be recorded and have explicit meaning
- *Example* - a student records database
  - Contains information identifying students, courses they are enrolled in, results from past courses, ...

Item	Type of data	Stored as
Family name	String	Character strings?
Birthdate	Date	3 integers?
Weight	Real number	Floating point number?
...		

# What is a database?

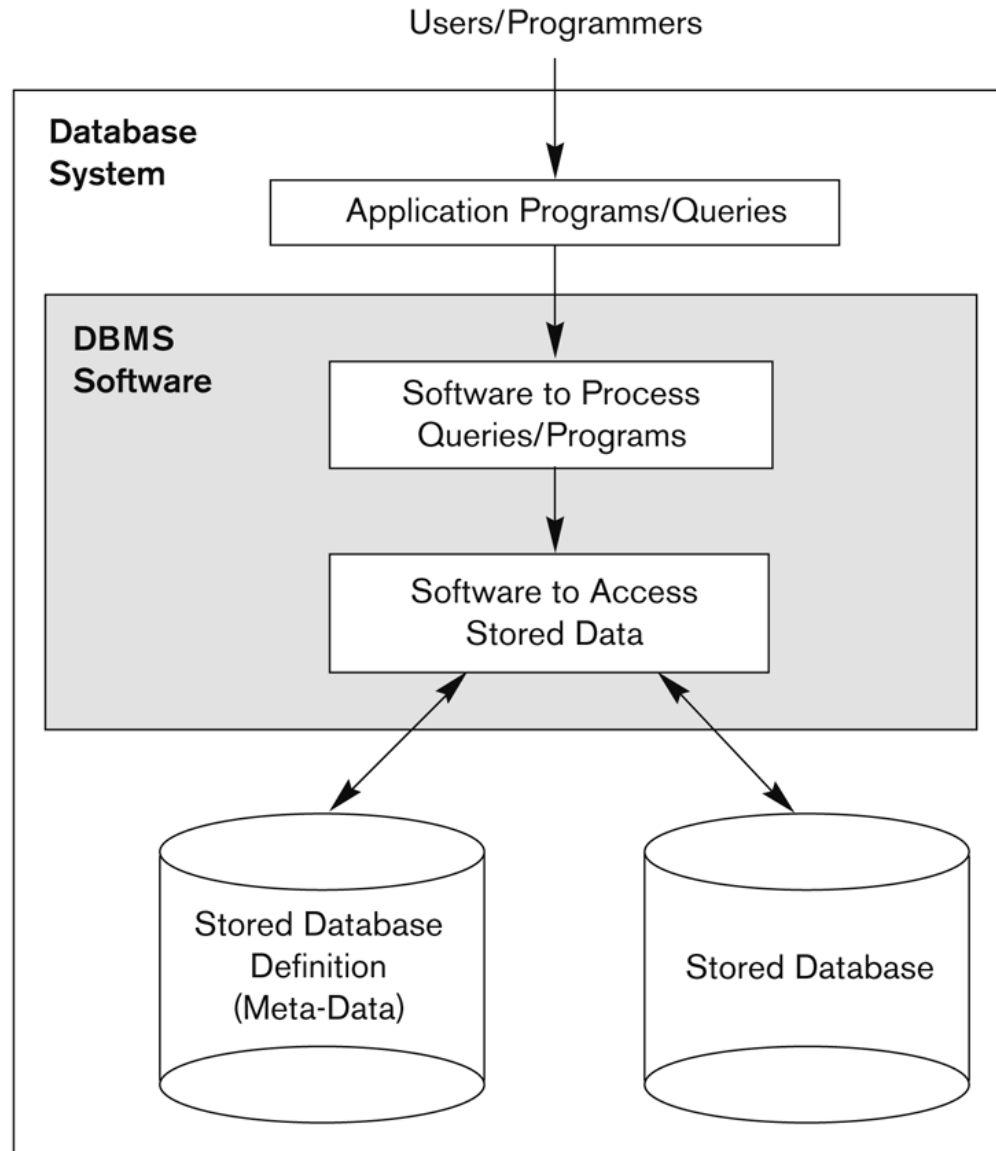
- *Database* (Elmasri/Navathe)
  - . . . a collection of **related** data . . .
- Data items alone are relatively useless
- We need the data to have some structure
- Database can be manipulated by a **database management system**

# What is a database management system (DBMS)?

- Elmasri/Navathe
  - *DBMS*: ... a collection of programs that enables users to create and maintain a database ...
  - *Database system*: ... the database and DBMS together ...



# Simplified Database System Environment (Elmasri/Navathe)



# Database Users

- *Database Administrator (DBA)*
  - Responsible for authorizing access to the database, coordinating and monitoring its use, acquiring software and hardware resources, controlling its use and monitoring efficiency of operations
- *Database Designer*
  - Responsible for defining the structure, the constraints, and transactions. Must communicate with the end-users and understand their needs

# Database Users (cont'd)

- *End Users* - Use the data for queries, reports and some of them update the database content. They can be categorized into:
  - **Casual**: access database occasionally when needed
  - **Naive** (or Parametric): they make up a large section of the end-user population.
    - They use previously well-defined functions in the form of “canned transactions” against the database.
    - Examples are bank-tellers or reservation clerks who do this activity for an entire shift of operations.

# Categories of End-Users (cont'd)

- **Sophisticated:** these include business analysts, scientists, engineers, others who are thoroughly familiar with system capabilities. Many use tools in the form of software packages that work closely with the stored database.
- **Stand-alone:** mostly maintain personal databases using ready-to-use packaged applications. An example is a user who maintains her/his own address book

# Database Management Systems

DBMS for practical work

- PostgreSQL 9.3.3 (open-source, free, full-featured)

Comments on using a specific DBMS:

- the primary goal is to learn SQL (a standard)
- the specific DBMS is not especially important
- but, each DBMS implements non-standard features
- we will use standard SQL as much as possible
- an exception is PL/pgSQL (but close to Oracle's PL/SQL)
- PostgreSQL documentations describe all deviations from standard

# Database Management Systems (cont'd)

## Comments on PostgreSQL vs Oracle

- Oracle is resource hungry (>800MB vs <200MB for PostgreSQL)
- PostgreSQL is a commercial-strength (ACID) RDBMS  
... but, being open source, you can see how it works
- PostgreSQL has been object-relational longer than Oracle  
... and its extensibility model is better than Oracle's
- PostgreSQL is more flexible than Oracle  
... allows stored procedures via a range of programming languages

But note: PostgreSQL and Oracle have very close SQL and PL/SQL languages

# Database Management Systems (cont'd)

## Comments on PostgreSQL vs MySQL

- both open source\* and reasonably efficient
- most Web/DB developers use MySQL
- until v4/5, MySQL lacked many serious DB concepts
  - no transactions, foreign keys, subselects, views, procedures, ...
- MySQL's SQL often ignores SQL standards
- MySQL is hacked together from "imported components"
  - multiple storage engines (some still w/o transactions)

PostgreSQL is better engineered; MySQL is more popular

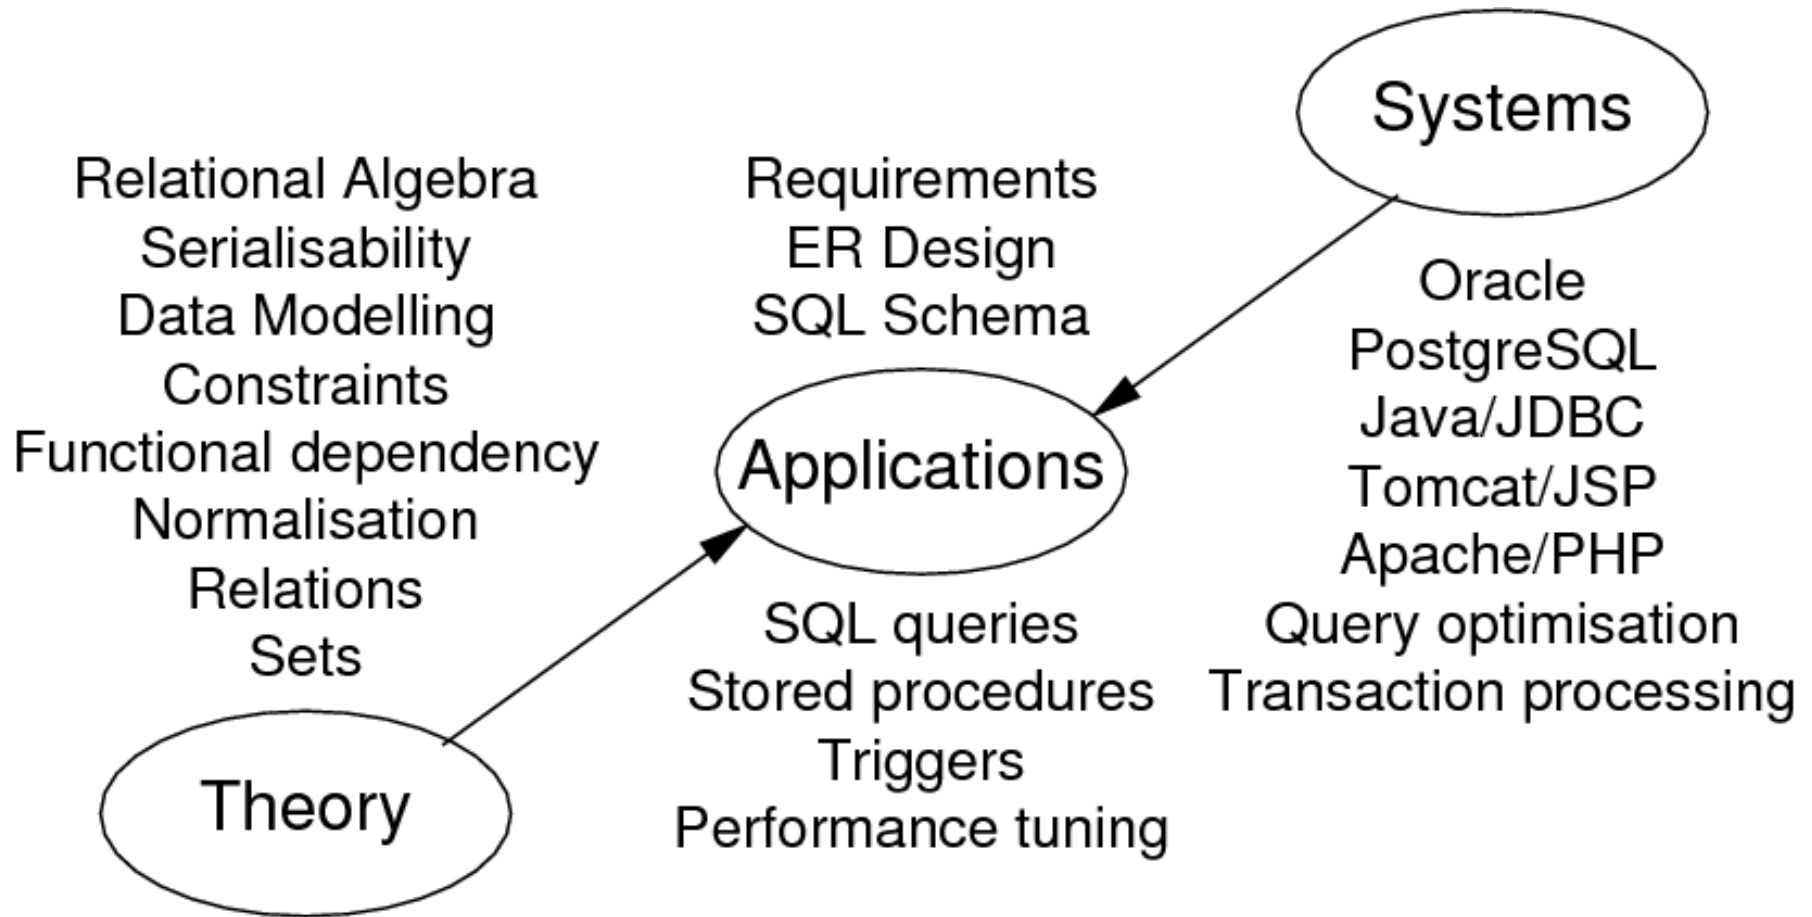
\* But Oracle now controls MySQL  $\Rightarrow$  open-source status unclear

# Working from Home

- PostgreSQL has very good on-line [documentation](#)
- This semester we will be using PostgreSQL 9.3.3
- If you install them at home
  - Get a version "close to" that (the latest stable is 9.6)
  - Test all assignment work at CSE before submitting
- Alternative to installing at home
  - Run them on CSE servers (grieg) as you would do in labs
  - Use, e.g., Putty to log in to a CSE server from home
  - PostgreSQL via Putty ok, since command-line based



# Overview of the Database Areas



# Database Application Development

A variation on standard software engineering process

1. analyse application requirements
2. develop a data model to meet these requirements
3. define operations (transactions) on this model
4. implement the data model as relational schema
5. implement transactions via SQL and PLs
6. construct a web interface to these transactions

At some point, populate the database (via interface?)

# Database System Languages

## Requests to DBMS

- **queries, updates** in data manipulation language (DML)  
(e.g. SQL)
- **data structures, constraints** in data definition language (DDL)  
(e.g. SQL)
- **create** and **drop** databases, **indexes, functions**  
(e.g. PL/pgSQL)

## Results/effects from DBMS requests

- **tuples** (typically, sets of tuples)
- changes to underlying data store

# Data Modeling

## Aims of data modeling

- describe what **information** is contained in the database  
(e.g., entities: students, courses, accounts, branches, patients, ...)
- describe **relationships** between data items  
(e.g., John is enrolled in COMP9311, Paul's account is held at Coogee)
- describe **constraints** on data  
(e.g. 7-digit IDs, students can enrol in no more than 30UC per semester)

## Data modelling is a **design** process

- converts requirements into a data model

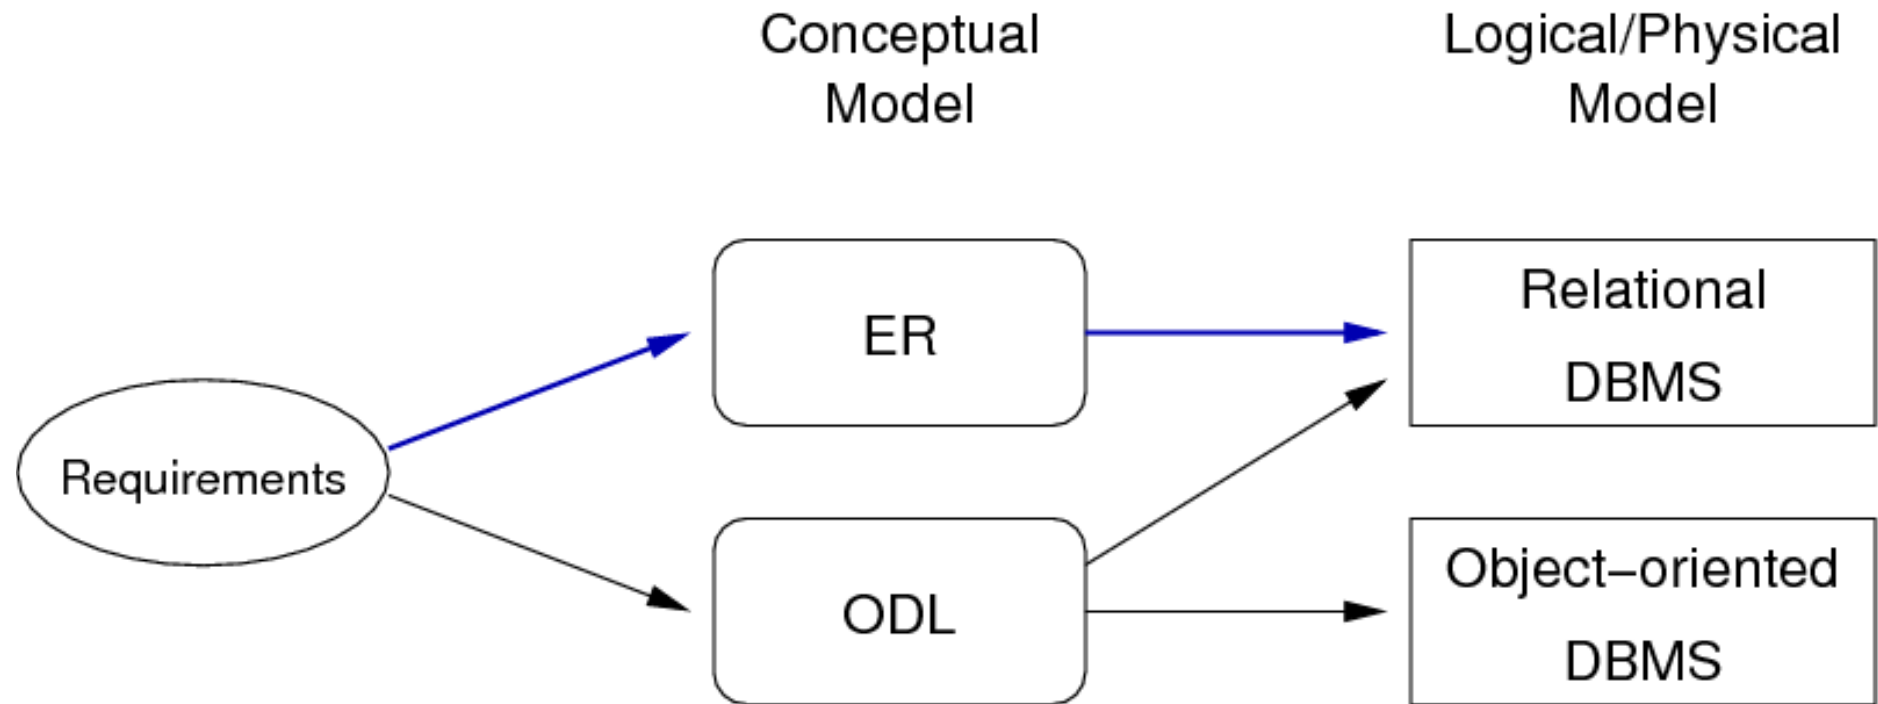
# Data Modeling (cont'd)

Kinds of data models:

- **logical**: abstract, for conceptual design  
e.g., Entity Relationship (ER), Object Definition Language (ODL)
- **physical**: record-based, for implementation  
e.g., relational

Strategy: design using abstract model; map to physical model

# Data Modeling (cont'd)



# Some Design Ideas

Consider the following while modeling data:

- start simple ... evolve design as problem better understood
- identify objects (and their properties), then relationships
- most designs involve kinds (classes) of people
- keywords in requirements suggest data/relationships  
(rule-of-thumb: nouns → data, verbs → relationships)
- don't confuse operations with relationships  
(operation: he buys a book; relationship: the book is owned by him)
- consider all possible data, not just what is available

# Exercise: Gmail Data Model

Consider the [Gmail](#) system

Develop an informal data model for it by identifying:

- the data items involved (objects and their attributes)
- relationships between these data items
- constraints on the data and relationships



# Quality of Designs

There is no single "best" design for a given application

Most important aspects of a design (data model)

- correctness (satisfies requirements accurately)
- completeness (all requirements covered, all assumptions explicit)
- consistency (no contradictory statements)

Potential **inadequacies** in a design

- omits information that needs to be included
- contains redundant information ( $\Rightarrow$  inconsistency)
- leads to an inefficient implementation
- violates syntactic or semantic rules of data model

# Entity-Relationship (ER) Model

The world is viewed as a collection of **inter-related entities**

ER has three major modelling constructs:

- **attribute**: **data item** describing a property of interest
- **entity**: collection of attributes describing **object** of interest
- **relationship**: **association** between entities (objects)

The ER model is not a standard. Hence, many variations exist

Lecture notes use notation from SKS and GUW books (simple)

# Entity-Relationship (ER) Diagrams

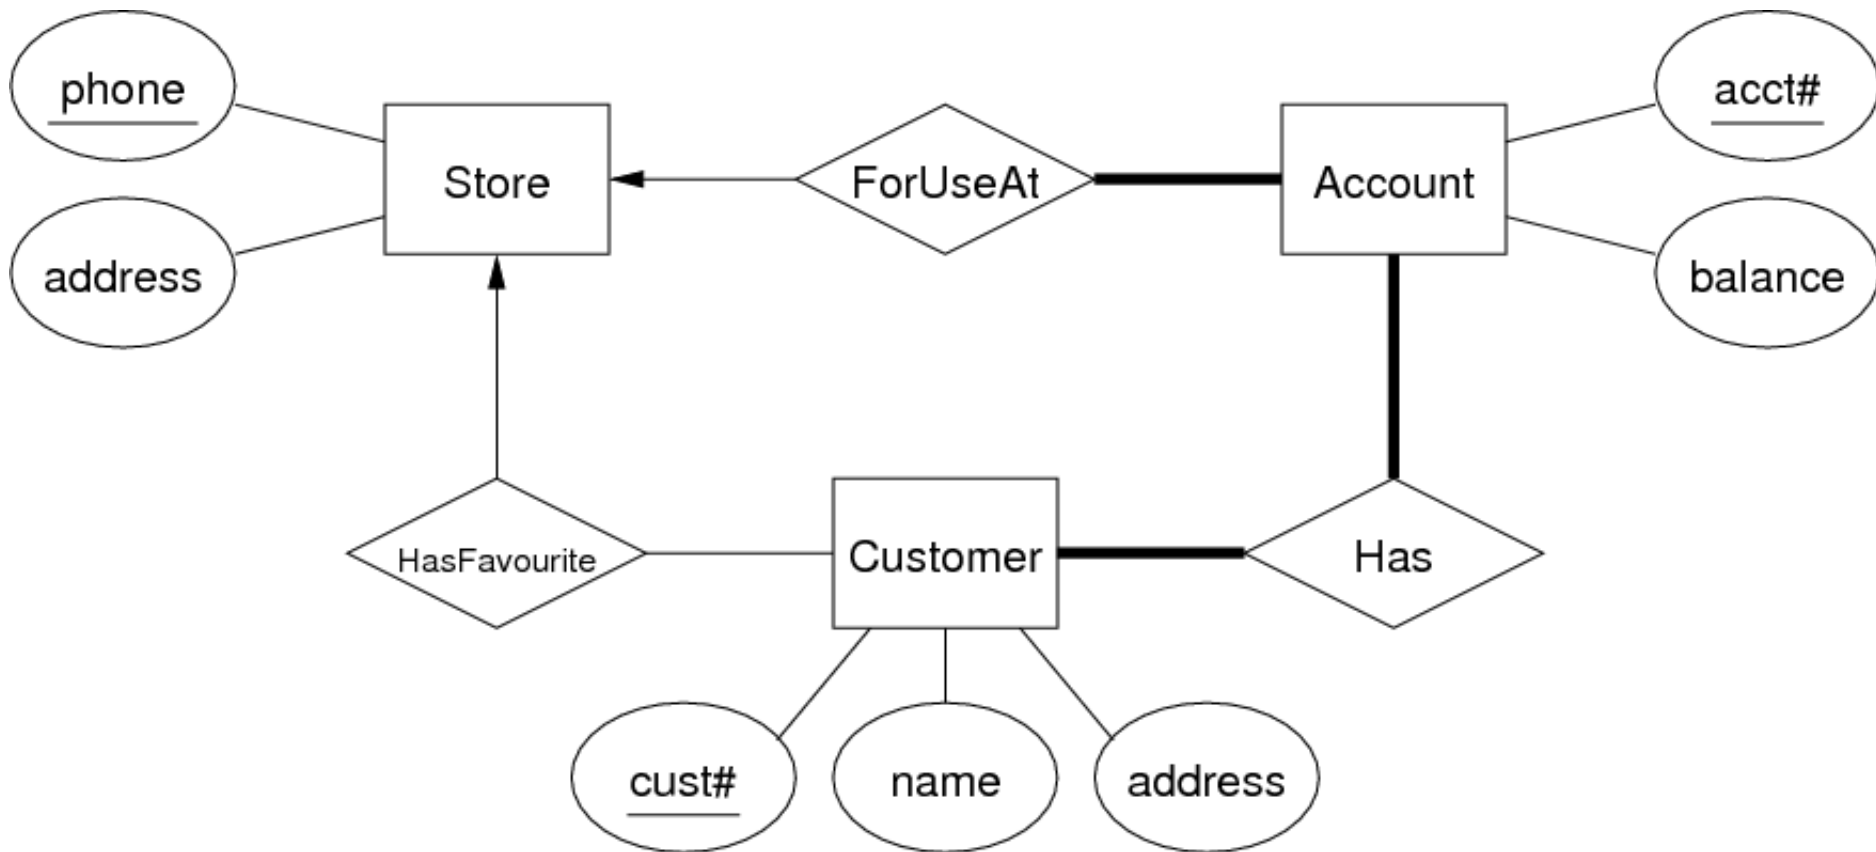
**ER diagrams** are a graphical tool for data modelling

An ER diagram consists of:

- a collection of **entity set** definitions
- a collection of **relationship set** definitions
- **attributes** associated with entity and relationship sets
- connections between entity and relationship sets

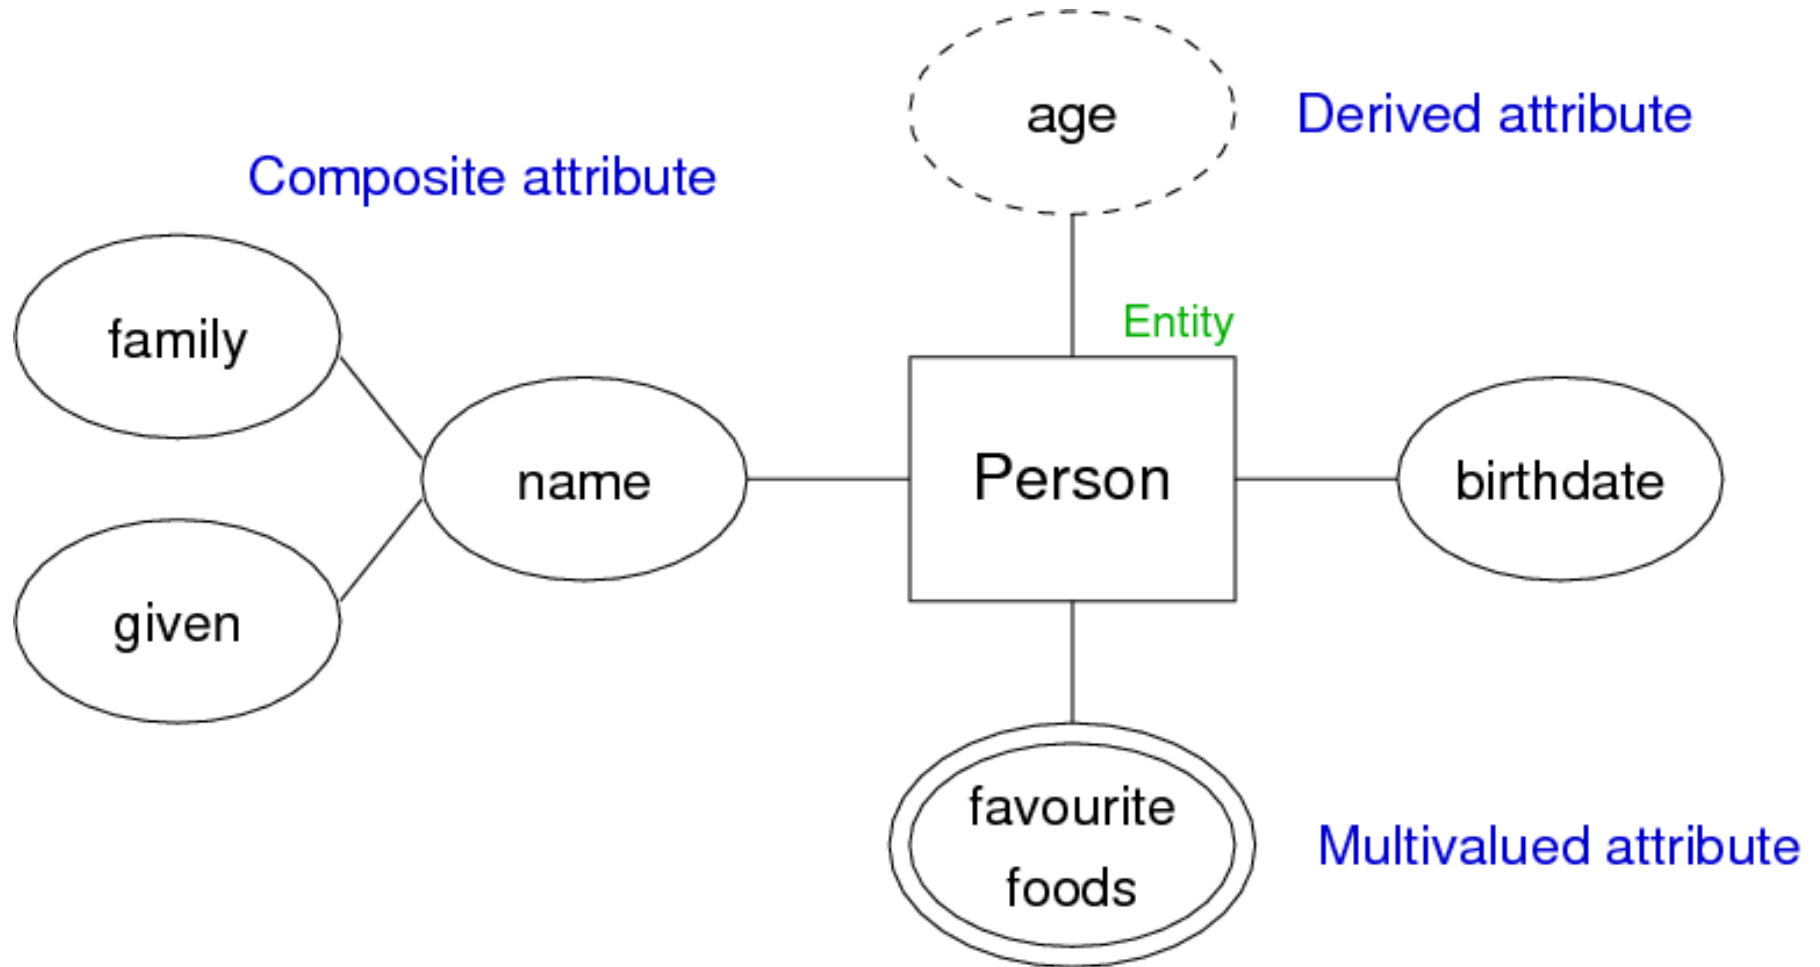
Terminology abuse: "entity" means "entity set" or "entity instance"?

# Entity-Relationship (ER) Diagrams (cont'd)



ER Diagram Example

# Entity-Relationship (ER) Diagrams (cont'd)



Attribute Notations Example

# Entity Sets

An **entity set** can be viewed as either:

- a set of entities with the same set of attributes (**extensional view** of entity set)
- an abstract description of a class of entities (**intentional view** of entity set)

An entity may belong to more than one entity sets

"Data" in a database  $\cong$  collection of (extensional) entity sets

# Keys

**Key (superkey)**: any set of attributes

- whose set of values are distinct over entity set
- natural (e.g., name+address+birthday) or artificial (e.g., SSN)

A **candidate key** is any superkey such that

- no proper subset of its attributes is also a superkey

A **primary key**:

- is one candidate key chosen by the database designer

Keys are indicated in ER diagrams by underlining

# Relationship Sets

**Relationship**: an association among several entities

E.g., Customer(9876) **is the owner** of Account(12345)

**Relationship set**: collection of relationships of the same type

**Degree** = # entities involved in relationship (in ER model,  $\geq 2$ )


**Cardinality** = # associated entities on each side of relationship

**Participation** = must every entity be in the relationship



# Relationship Sets (cont'd)

**Example:** relationship participation




# Relationship Sets (cont'd)

**Example:** relationship cardinality


one-to-one



one-to-many

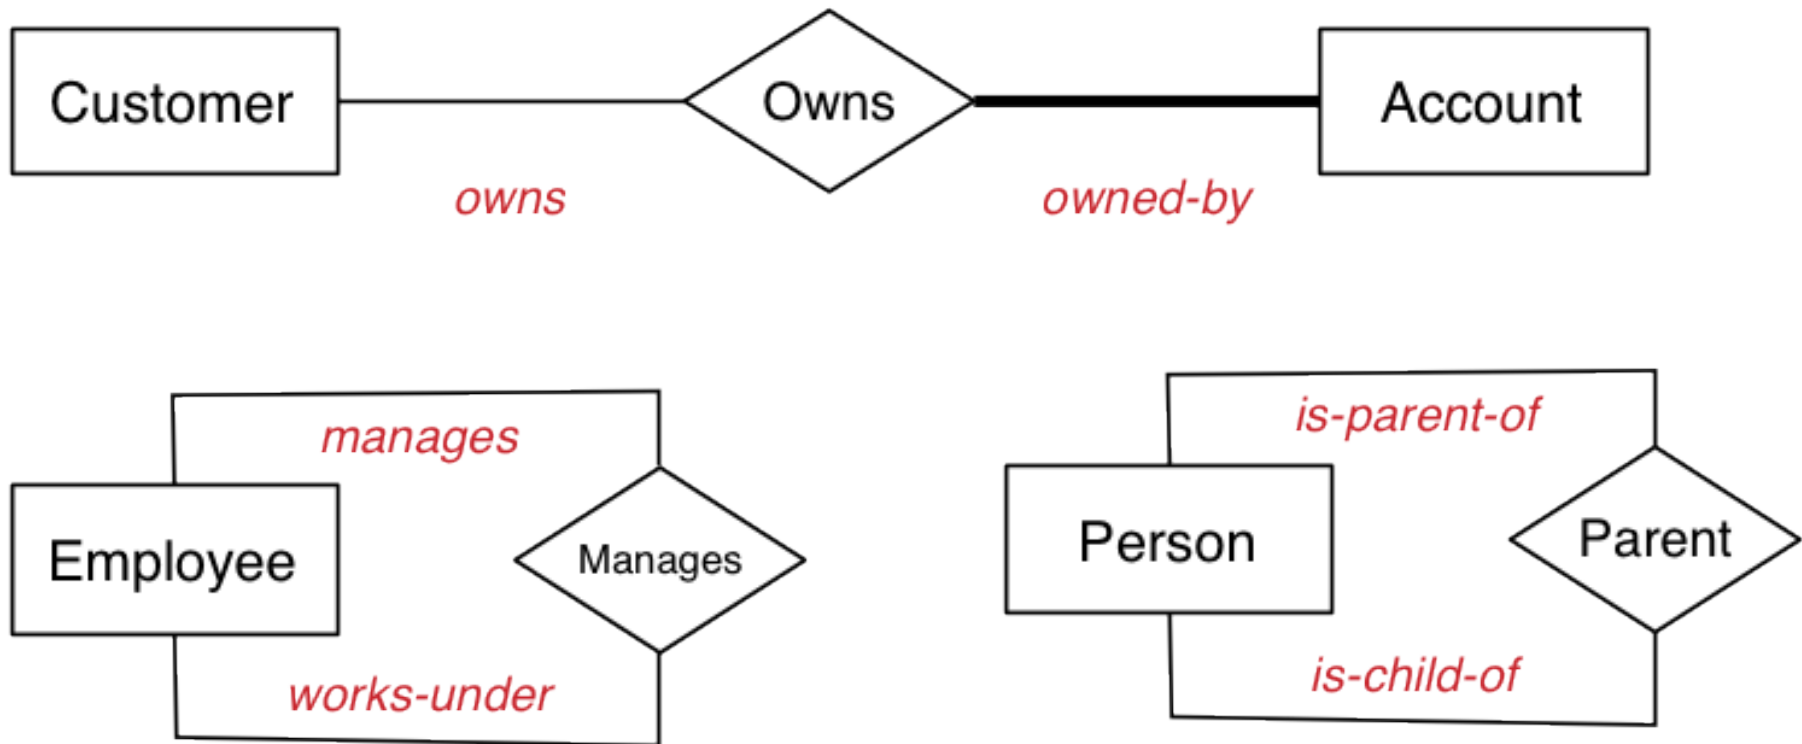


many-to-many



# Relationship Sets (cont'd)

The role of each entity in a relationship is usually implicit.  
If ambiguity arises, can explicitly name the role.




Role names become more important when developing SQL schemas

# Relationship Sets (cont'd)

In some cases, a relationship needs associated attributes.

Example:




(Price and quantity are related to products in a particular shop)

# Weak Entity Sets

## Weak entities

- exist only because of association with strong entities
- have no key of their own; have a **discriminator**

Example:



# Subclasses and Inheritance

A **subclass** of an entity set  $A$  is a set of entities:

- with all attributes of  $A$ , plus (usually) its own attributes
- that is involved in all of  $A$ 's relationships, plus its own

Properties of subclasses:

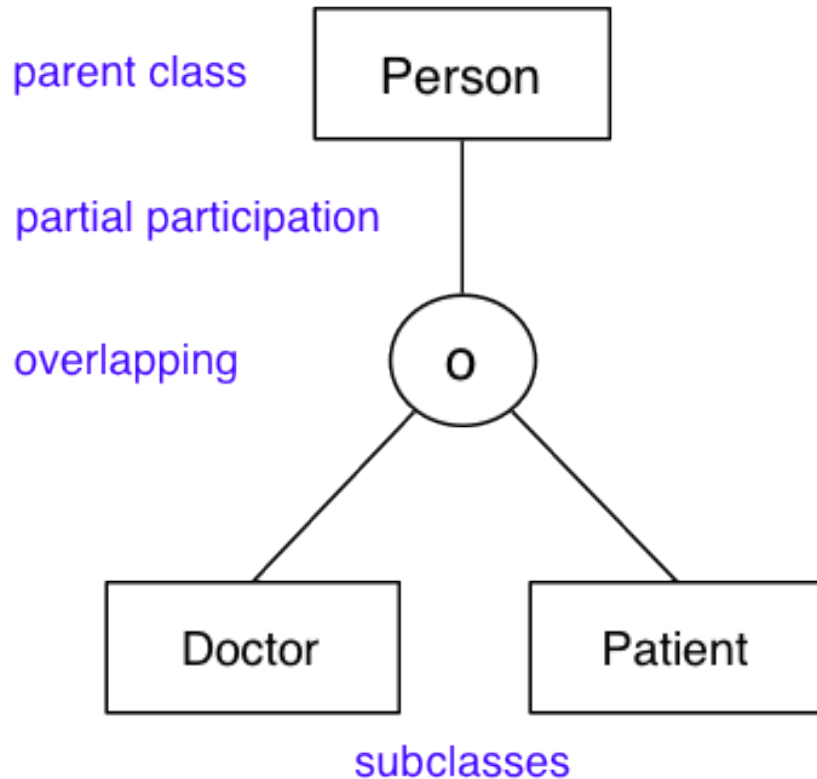
- **overlapping** or **disjoint** (can an entity be in multiple subclasses?)
- **total** or **partial** (does every entity have to also be in a subclass?)

Special case: entity has one subclass (" $B$  **is-a**  $A$ " specialisation)


# Subclasses and Inheritance (cont'd)

Example:

*A person may be a doctor and/or may be a patient or may be neither*



*Every employee is either a permanent employee or works under a contract*



# Design Using the ER Model

ER model: simple, powerful set of data modelling tools

Some considerations in designing ER models:

- should an "object" be represented by an attribute or entity?
- is a "concept" best expressed as an entity or relationship?
- should we use  $n$ -way relationship or several 2-way relationships?
- is an "object" a strong or weak entity? (usually strong)
- are there subclasses/superclasses within the entities?

Answers to above are worked out by *thinking* about the application domain



# Design Using the ER Model (cont'd)

ER diagrams are typically too large to fit on a single screen (or a single sheet of paper, if printing)

One commonly used strategy:

- define entity sets separately, showing attributes
- combine entities and relationships on a single diagram (but without entity attributes)
- if very large design, may use several linked diagrams

# Exercise: Medical Information

- Patients are identified by an SSN, and their names, addresses and ages must be recorded.
- Doctors are identified by an SSN. For each doctor, the name, specialty and years of experience must be recorded.
- Each pharmacy has a name, address and phone number. A pharmacy must have a manager.
- A pharmacist is identified by an SSN, he/she can only work for one pharmacy. For each pharmacist, the name, qualification must be recorded.
- For each drug, the trade name and formula must be recorded.
- Every patient has a primary physician. Every doctor has at least one patient.
- Each pharmacy sells several drugs, and has a price for each. A drug could be sold at several pharmacies, and the price could vary between pharmacies.
- Doctors prescribe drugs for patients. A doctor could prescribe one or more drugs for several patients, and a patient could obtain prescriptions from several doctors. Each prescription has a date and quantity associated with it.